

# Assignment # 1: A Simple Shell<sup>1</sup>

Version 0.2 – 10 July 2014 – draft only

## COMP3301/COMP7308: Operating Systems Architecture

Due Friday 22<sup>nd</sup> August, 2014, 8pm. (End of week 4)

This assignment asks you to write your own shell, called **itsh** (ITEE Shell), using the fork, exec, wait, dup2, and pipe system calls. The purpose is to get you familiar with Unix processes, creating them, passing arguments, inheriting open files, waiting for them to terminate, background processes, and so on. You will also be introduced to pipes as well. You will write your shell in C.

Submission will be via a gzipped tarball that includes a makefile.

### 1 Overview

In this assignment you will write your own shell, called **itsh**. This will work like the Bash shell you are used to using, prompting for a command line and running commands, but it will not have many of the special features of the Bash shell.

- Your shell will support:
  - the redirection of standard input
  - the redirection of standard output
  - use of | (an unnamed pipe)
  - foreground processes
  - background processes
  - comments (lines beginning with the # character).
- and the following two built in commands:
  - exit
  - cd

### 2 Specifications

Here are the specifications for your assignment. There should be no guess work on what is required for this assignment. If you have any questions, check the Discussion Board on Blackboard. If your questions aren't answered there, post your question there.

---

<sup>1</sup> This assignment is based on a similar assignment offered in the course CS311 at Oregon State University, USA: [classes.engr.oregonstate.edu/eecs/winter2011/cs311/Winter11Projects/Proj4/assign4.doc](http://classes.engr.oregonstate.edu/eecs/winter2011/cs311/Winter11Projects/Proj4/assign4.doc) and on many other similar “write a shell” assignments.

## 2.1 Command Line

- The general syntax of a command line is

```
command [arg1 arg2 ...] [< input_file] [> output_file] [&]  
or  
command [arg1 arg2 ...] [> output_file] [< input_file] [&]  
or  
command [arg1 ...] [< input_file] | command [arg1 ...] [> output_file]
```

where items in square brackets are optional.

- You can assume that a command is made up of arguments (including characters such as > < & or |) separated by one or more whitespace characters. There may be whitespace characters before the first command word. For testing, only space characters will be used for whitespace.
- The special symbols |, <, >, and & are recognized, but they must be surrounded by spaces like other words.
- You can assume that < will appear at most 1 time in the command. You can assume the > will appear at most 1 time in the command. You can assume that | will appear at most 1 time in the command. However, they may appear together in the same command. For example,  
cat myFile | wc > otherFile  
wc -l < myFile > otherFile  
sort < myFile | uniq > otherFile
- The command on the left side of a pipe will not also have output redirection.
- The command on the right side of a pipe will not also have input redirection.
- If the command is to be executed in the background, the last word must be &. As the above general syntax indicates, you do not have to worry about piped commands running in the background.
- If standard input or output is to be redirected, the < or > words followed by a filename word must appear after all the arguments (i.e., the command and arguments occur before any I/O file redirection.)
- Input redirection can appear before or after output redirection.
- Your shell does not need to support any quoting; so arguments with spaces inside them are not required to be supported.
- Valid command lines may be up to 128 characters, including newline, with at most 20 arguments or special characters. This includes both commands if the command line includes a pipe.
- You may choose to execute lines longer than 128 characters if your data structures support them, or you may print an error message.
- Your program should not crash with arbitrarily long lines.
- Use the current working directory followed by a hash (#) as a prompt for each command line.

## 2.2 Command Execution

- For executing commands use:
  - `fork()`  
to start new processes
  - `dup2()`  
to duplicate file descriptors
  - `pipe()`  
to create a unnamed pipe
  - `exec()` family  
to execute new programs
  - `waitpid()`  
You will need to use `waitpid` to check for completed background processes. Read about the `WNOHANG` option of `waitpid`. It might be useful.
  - Your shell should use the `PATH` variable to look for commands. The right version of the `exec` function will do this for you automatically.
  - Your shell should allow shell scripts to be executed. In other words, `itsh` should take an optional argument which is the name of a file from which to read commands.
- If a command fails because the shell could not find the command to run, then the shell will print an error message. You can use macros/functions such as `errno`, `perror` or `strerror` to generate appropriate error messages on `stderr`.
- If the files for file redirection cannot be opened, then the command should not be invoked.
- If one of the commands in a pipe fails then either the whole command should fail, or the correct command may continue to run, perhaps terminating because it requires I/O from the pipe. Any of these behaviours is OK – whatever is easiest to implement.
- Background and foreground commands
  - Your shell will wait for the completion of foreground commands (commands without the `&` in the final position) before prompting for a new command.
  - Your shell will NOT wait for commands run in the background. Instead it will prompt for a new command without waiting for background commands to complete.
  - Background commands should have their standard input redirected to be from `/dev/null` if the user did not specify some other file to take standard input from.
- Input redirection

After the `fork` but before the `exec` you must do any input/output redirection.

  - A redirected input file should be opened for reading only
  - If your shell cannot open the file for reading it should print an error message and the command will not run.
- Output redirection

After the `fork` but before the `exec` you must do any input/output redirection.

  - A redirected output file should be opened for write only,
  - The output file should be truncated if it already exists or created if it does not exist.

- If your shell cannot open the output file it should print an error message and the command will not run.
- Piping
  - Create the pipe before the fork, then do input / output redirection, then exec.
  - A redirected input file should be opened for reading only
  - If your shell cannot open the file for reading it should print an error message and the command reading from it should fail.
  - A redirected output file should be opened for write only,
  - It should be truncated if it already exists or created if it does not exist.
  - If your shell cannot open the output file it should print an error message and the command writing to it should fail should fail.

## 2.3 General Behavior

- control-C interrupt from the keyboard does NOT terminate your shell,
- control-C interrupt from the keyboard does terminate the foreground command currently running.
- Background commands should not be terminated by a control-C signal.
- Your shell should allow blank lines and comments.
- Any line that begins with the # character is a comment line and should be ignored.
  - #This is a comment.
  - # Another comment.
- For comments greater than 128 characters you either ignore the whole line, or else you can indicate an error.
- A blank line (one without any commands) should do nothing; your shell should just output a prompt for another command (like Bash does).
- The shell will print the process id of a background process when it begins.
- When a background process terminates, a message showing the process id and exit status will be printed.
- You should check to see if any background processes completed just before you prompt for a new command and print the message then. Note: In this way the messages about completed background processes will not appear during other running commands, though the user will have to wait until they complete some other command to see these messages (this is the way the C shell and Bourne shells work).

## 2.4 Built in Commands

- Your shell will support two built in commands: exit and cd.
- You do not have to support input/output redirection, running in the background, or pipes for these built in commands.
- Built-in command specifications
  - *exit*
    - `exit`
    - The exit command exits the shell.
    - It takes no arguments.
    - If arguments are supplied, then they are ignored.

- *cd*

```
cd [dir]
```

- Purpose: The `cd` command changes directories.
- when called with no arguments, it changes to the directory specified in the *HOME* environment variable.

You can get your process' value of *HOME* with `getenv("HOME")`.

You can change directories with `chdir()`.

- It can also take one argument, the path of the directory to change to.
- If it is passed more than one argument, the extra arguments are ignored.

For example:

```
cd dir1 dir2
```

Will have the same effect as:

```
cd dir1
```

## 2.5 ‘man’ page

As documentation for your **itsh**, you need to produce a “man” page for your shell, with a maximum of 10 kbytes of characters. Any lack of functionality can be explained in the “bugs” section.

Format your man page in accordance with the conventions for writing Linux man pages available at

```
man man-pages
```

Include at least the following sections:

NAME

SYNOPSIS

DESCRIPTION

OPTIONS [Normally only in Sections 1, 8]

EXIT STATUS [Normally only in Sections 1, 8]

ENVIRONMENT

FILES

NOTES

BUGS

EXAMPLE

SEE ALSO

### 3 Submission

The shell should be submitted as a gzipped tarball, <studentnumber>.tar.gz  
It should include a makefile which produces a “itsh” executable in the current working directory. You should “make clean” before producing the tarball.

The makefile should implement the following targets:

all (which produces the executable in the working directory and copies the man page if necessary)

clean (removes the executable and object files)

dist (generate the gzipped tarball after first doing a make clean)

You can use whatever compiler options you wish, but you should only use standard C libraries unless otherwise allowed on the discussion board. C99 extensions are allowed.

Currently allowed extensions include:

- Use of lex and yacc or flex and bison programs to generate input line parsers (Note: If you use these you should submit your lex/yacc or flex/bison source code and call these programs as part of your makefile.)
- Use of GNU readline library to handle input lines from user.

The tarball should include a man page in the current working directory called itsh.1

You can submit your tarball file via the “Assessment” link on Blackboard.

You may use a version control system, such as subversion, to develop your program, but submission is only of the finished product as a tarball.

You should confirm that your code runs on *moss.labs.eait.uq.edu.au*, because this is where it will be tested.

### 4 Marking Scheme

- WARNING: If your code doesn’t work, you won’t get credit for it.
  1. For this assignment, break the assignment into pieces that you can work on independently.
  2. Get one piece fully functional and thoroughly tested.
  3. Then go on to the next piece.
  4. If you’ve tried to implement everything but haven’t completed any pieces, you won’t get credit for anything.

Individual functions will be checked, tested and marked.  
For each function, percentages of marks will be allocated as  
0-20% Function not implemented, or no working code.  
20-40% Some simple but incorrect functionality  
40-60% Significant problems, such as frequent crashes or infinite loops  
60-80% Moderate problems, fails relatively often.  
80-100% Few or no problems.

Total marks out of 100:

Makefile works, program which addresses the specifications compiles [10]  
Coding style, readability, organization [10]  
Foreground Execution [20]  
Background Execution [10]  
Built-in Command exit [5]  
Built-in Command cd [5]  
Input redirection [10]  
Output redirection [10]  
Pipes [10]  
Man page [10]

## ***5 Test File***

Some example testfiles are provided on Blackboard.  
One called “longtest” is a here file that sort of mimics interactive input.  
The second “filetest” is the same set of commands but as an input file for itsh.

## ***6 Revision History***

Version 0.1, 7/7/2014 – Draft for Comment.  
Version 0.2, 10/7/2014 – flex & bison added as permitted packages